

Zero of twenty-six, and that is the expected number

A mechanical check of the working tree against the Target Architecture Specification v1.0. No requirement is met yet — but the reason matters more than the number: nothing has been attempted, and almost everything the specification depends on is already in the repository.

0/26

Specification requirements satisfied. The working tree is byte-identical to the tree that was audited: **no source file has changed**, and no commit has landed since `1f1cbbb`.

26 GAPS • 0 MET • 0 IN PROGRESS

Project

Anamnesis Media Forensics Engine

Measured

31 August 2026

Tree

1f1cbbb + 0 source changes

Method

Scripted, re-runnable

READ THIS NUMBER CORRECTLY

0/26 measures **conformance, not readiness**. The audit was analysis-only by instruction and no implementation was authorised, so a zero was the only honest result available today. Its value is as a baseline: every future run of the same script is directly comparable to this one.

What the number does *not* say is that the work is twenty-six units away. Twelve of the twenty-six are single-line or single-file changes, and the three largest all consume code that already exists and already compiles.

26 requirements open	1,088 lines to reconnect, not write	3,130 lines to delete outright	3 genuinely new files needed	~400 net new lines, all told
--------------------------------	---	--	--	--

The audit found a codebase whose architecture was sound and whose content was fabricated. This report measures the distance between that codebase and the specification issued to correct it. The distance is real, but it is almost entirely made of deletions and reconnections rather than authorship — which is the same finding that produced the improve-rather-than-rebuild verdict, now expressed as a work estimate.

What was checked

Twenty-six assertions, each evaluated against the working tree by script rather than by reading: import-graph reachability from `src/main.tsx`, pattern checks for the specific constructs each requirement names, and dependency-manifest inspection. The row identifiers are stable, so progress from here is a diff rather than a re-read.

02 The ledger

Each row is a checkable assertion. The evidence column reports what the check actually found in the tree — not a restatement of the requirement.

ID	REQUIREMENT	STATE	EVIDENCE IN THE TREE	PHASE
G11	Model ID from environment	GAP	2 hardcoded <code>gemini-3.7-flash</code> literals	1
G2	Fabricated fallback report removed	GAP	3 Sulawesi / 2018 literals in <code>App.tsx</code>	1
G1	Client calls <code>/api/forensics/analyze</code>	GAP	no such fetch anywhere in <code>src/</code>	1
G22	Step 1 base64 consumed by <code>App</code>	GAP	handler declares 1 param; arg 2 dropped	1
G9	<code>lib/api.ts</code> typed fetch wrapper	GAP	file does not exist	1
G7	<code>res.ok</code> checked before rendering	GAP	response body parsed unconditionally	1
G8	Timeout / <code>AbortController</code>	GAP	absent on every client fetch	1
G18	Media stripped before <code>localStorage</code>	GAP	full data URLs persisted	1
G3	<code>buildCaseState</code> is a pure map	GAP	ENF, 12.4s, 14 copies, Palu still literal	2
G4	<code>NOT_ASSESSED</code> sentinel exists	GAP	not defined in <code>types.ts</code>	2
G20	No unreachable modules under <code>src/</code>	GAP	18 files unreachable from <code>main.tsx</code>	2
G21	<code>recharts</code> removed	GAP	declared; imported in 0 files	2
G10	<code>lib/statusStyles.ts</code> single badge map	GAP	4 components define <code>getStatusBadge</code>	2
G26	Product name spelled consistently	GAP	"ANAMESIS" in chat greeting + PDF header	2
G5	<code>ForensicCanvasLab</code> on live path	GAP	unreachable	3
G17	<code>contentType</code> sent with image bytes	GAP	omitted; server defaults to <code>image/jpeg</code>	3
G6	<code>ForensicCopilotChat</code> on live path	GAP	unreachable	3
G24	Hash race fixed (readers awaited)	GAP	VERIFIED shows before readers settle	3
G23	Upload size / type guard	GAP	UI claims 100 MB; nothing enforces	3
G12	<code>PORT</code> from environment	GAP	<code>const PORT = 3000</code>	4
G13	Production is the default branch	GAP	still <code>!= 'production'</code>	4

ID	REQUIREMENT	STATE	EVIDENCE IN THE TREE	PHASE
G19	Dockerfile present	GAP	absent	4
G14	Shared-secret auth on POST routes	GAP	both routes open to the internet	4
G15	Rate limiting	GAP	none	4
G16	Raw model output not echoed	GAP	raw: rawText returned on parse failure	4
G25	PDF emoji stripped	GAP	emoji written into jsPDF core fonts	4

03 What is already in place

The ledger measures behaviour. This table measures inventory — the assets the specification depends on that already exist, already compile, and need no authoring. This is the half of the picture a conformance percentage cannot show.

ASSET	LINES	STATUS	WHAT IT MEANS FOR THE PLAN
<code>src/types.ts</code>	271	READY	Both contracts already modelled. G4 adds a sentinel; nothing else changes.
<code>responseSchema</code> in <code>server.ts</code>	~130	READY	Returns exactly the type <code>buildCaseState</code> accepts. G1 is a wire, not an integration.
<code>lib/imageForensics.ts</code>	250	ORPHANED	Real ELA / noise / Sobel, already correct. G5 mounts it.
<code>ForensicCanvasLab.tsx</code>	619	ORPHANED	The UI for the above. Reconnected, not rewritten.
<code>ForensicCopilotChat.tsx</code>	219	ORPHANED	Needs G7 and G17 before it is safe to mount; then G6 mounts it.
<code>lib/cryptoUtils.ts</code>	136	LIVE	Correct SHA-256 and EXIF parser. Untouched by the plan.
<code>lib/soundFx.ts</code>	482	LIVE	Untouched by the plan.
Step 1–5 components and shell UI	~2,900	LIVE	Read the view model already. Behaviour changes upstream of them, not inside them.
Build pipeline	—	LIVE	<code>tsc --noEmit</code> clean; <code>vite build</code> in 1.8s. No build work in any phase.

THE ASYMMETRY THAT DEFINES THIS PROJECT

1,088 lines get reconnected. 3,130 lines get deleted. Roughly 400 get written. The largest single category of work in the entire plan is removal — and the second largest is re-importing files that already compile. That is why the verdict was improve rather than rebuild, and it is why a zero-percent conformance score is not alarming.

04 Progress by phase

PHASE 0	Decide and specify	3 / 3
Architectural audit complete; verdict reached (improve, 85% confidence); Target Architecture Specification v1.0 issued. The only phase that is finished — and the one that makes the rest cheap, because nobody now has to rediscover why the constants are being removed.		
PHASE 1	Critical stabilisation — stop the app from lying	0 / 8
G11 · G2 · G1 · G22 · G9 · G7 · G8 · G18. Ends when an unfamiliar photograph produces either a real analysis or a visible error, and never a stored report. Everything else is optional relative to this.		
PHASE 2	Architectural cleanup — drain the constants	0 / 6
G3 · G4 · G20 · G21 · G10 · G26. Mostly deletion. Ends when two different cases opened side by side show different findings.		
PHASE 3	Feature readiness — real capability back on screen	0 / 5
G5 · G17 · G6 · G24 · G23, plus Demo Mode. Ends when Step 2 renders an ELA computed in the browser from the ingested image.		
PHASE 4	Deployment and demo hardening	0 / 7
G12 · G13 · G19 · G14 · G15 · G16 · G25, plus two rehearsals on conference wifi with an unfamiliar file.		

Phases are ordered by dependency, not by size. Phase 1 is the largest and must go first: Phase 2 cannot safely delete the invented constants until something real is arriving to replace them, and Phase 3 must not mount the chat until Phase 1's error handling exists — otherwise it ships a component that answers every question with a fabricated all-clear.

05 Critical path from here

If only one day is available, these six steps in this order. Each is independently shippable and each leaves the tree in a working state.

1 Verify `gemini-3.7-flash` resolves — G11

One `curl` against the models endpoint. It gates the entire plan: if the model ID is wrong, Phase 1 cannot complete and the right response is a scope decision, not more code. Nothing else should start before this answer exists.

2 Delete the fabricated fallback — G2

Replace the 80-line literal at `App.tsx:319` with an explicit "no analysis available" state. Ship this even if nothing else lands today; it is the difference between an incomplete demo and a dishonest one.

3 Add `lib/api.ts` — G9, G7, G8

One typed wrapper that checks `res.ok`, aborts at 30 seconds and normalises errors. Three requirements close together, and no component can subsequently make the mistake individually.

4 Wire the analyze call — G1, G22

Restore the dropped second parameter, POST through the new wrapper, feed the response to `buildCaseState`. The endpoint and the type already match.

5 Strip media before persisting — G18

Prevents silent case loss the moment a real upload enters the flow — which step 4 makes possible for the first time.

6 Re-run this ledger

Expect 8 / 26. That is Phase 1 complete and the application no longer capable of fabricating a finding.

THE ONE SEQUENCING RULE

Do not mount `ForensicCopilotChat` (G6) before G7 and G17 are closed. Mounted as it stands, every benchmark case sends an SVG labelled as JPEG, the model rejects it, and the unchecked response renders as *"Analysis completed with no additional discrepancies found."* A chat that confidently clears every question is worse than no chat at all.

06 Honest risk position

RISK	STATUS	NOTE
Demo fabricates a finding on an unfamiliar file	LIVE	Unchanged since the audit. Highest-consequence open risk; closed by G2 alone.
Model ID may not resolve	UNVERIFIED	Still unchecked. Cheapest check in the plan, largest blast radius.
Public unauthenticated Gemini proxy	LIVE	Deployed URL is in the README. Quota exhaustion before judging is a realistic outcome.
Silent case loss on real uploads	LIVE	Currently masked because benchmark SVGs are small. G1 makes real uploads routine and unmask it.
Scripted benchmark demo breaks	LOW	Deterministic and offline by construction. Unaffected by any phase — Demo Mode preserves it deliberately.
Build or typecheck regression	LOW	Both clean at baseline and re-checkable in seconds.

Nothing has got worse since the audit, and nothing has got better. The risk position is exactly the one recorded earlier today — which is the correct outcome of an analysis-only pass, and is the last report in which that will be an acceptable answer.

Next measurement

Re-run the ledger at the end of each phase. Expected trajectory: **8/26** after Phase 1, **14/26** after Phase 2, **19/26** after Phase 3, **26/26** after Phase 4. A phase that closes fewer rows than expected is the signal to re-examine the plan rather than to press on.

Measured by scripted assertion against the working tree at commit 1f1cbbb with zero uncommitted source changes: import-graph reachability from src/main.tsx, pattern checks per requirement, and dependency-manifest inspection. Untracked files present: package-lock.json, ANAMNESIS-Target-Architecture.pdf. No source file was modified to produce this report.